



Probability and Statistics for Business and Data

PART 6 - REGRESSION




Linear Regression



Linear Regression

- The goal of **regression** is to develop an equation or formula that **best describes** the relationship between variables.




$$y = 2x$$



Linear Regression

- How do we find a best-fit line?
- Consider a dataset with only one variable
- The best-fit line is just the mean value of the data points





Understanding Best Fit


- A plant manager wants to know when employees arrive at work
- The shift starts at 8:30am
- She takes five random timecards and plots the minutes of arrival on a chart





Understanding Best Fit

Timecard	Minutes past 8:00am
A	37
B	25
C	27
D	36
E	25
Total:	150
Mean	30

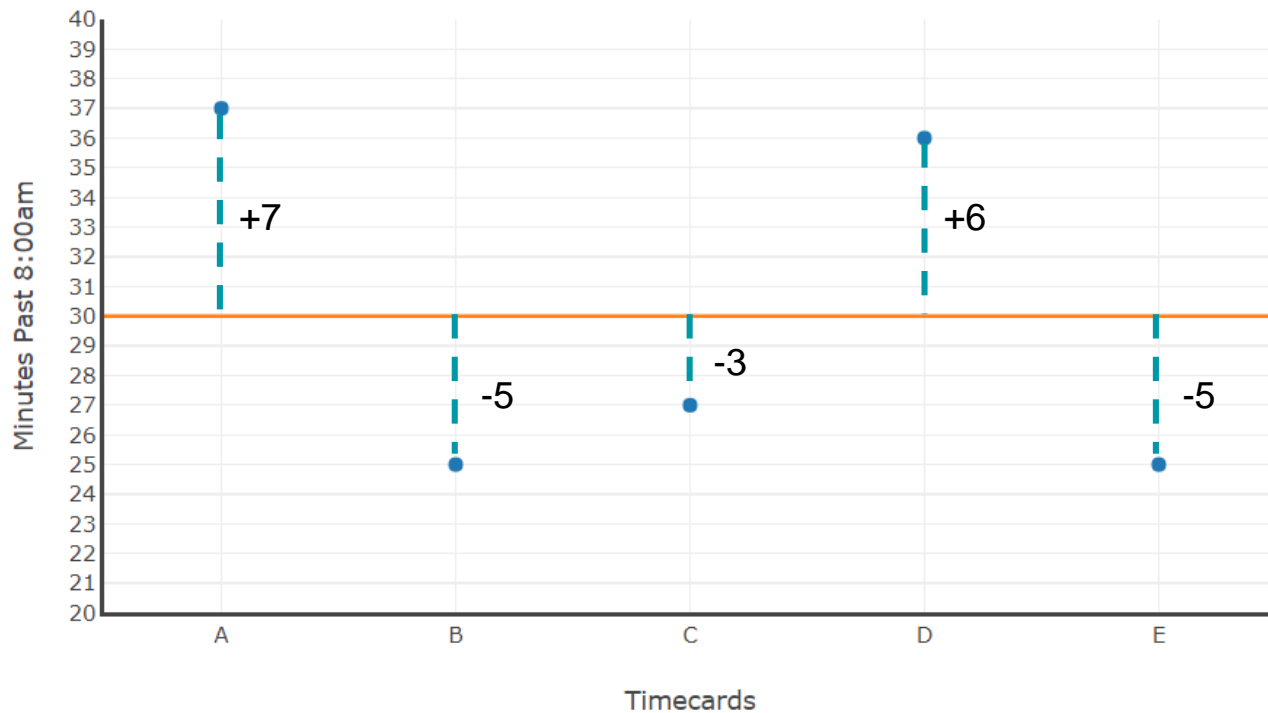




Understanding Best Fit

What makes
 $y = 30$ a
best-fit line?


Consider the
error





Understanding Best Fit


See that the sum of the distances above the line balances the sum of those below the line





Understanding Best Fit

Error (E)	Square Error (SE)
+7	49
-5	25
-3	9
+6	36
-5	25
Sum of Squares Error (SSE)	144



we want to MINIMIZE the SSE




Understanding Best Fit

What if we
move the line?

Let's set it to
 $y = 31$ instead


How does it
affect the SSE?





Understanding Best Fit

Error (E)		Square Error (SE)	
+7	+6	49	36
-5	-6	25	36
-3	-4	9	16
+6	+5	36	25
-5	-6	25	36
Sum of Squares Error (SSE)		144	149




moving the line INCREASED the SSE



Linear Regression




- That's it! The goal of regression is to find the line that best describes our data.
- Fortunately, we don't have to rely on trial-and-error.
- We have algebra!





Linear Regression

- Recall that the equation of a line follows the form $y = mx + b$ where
 m is the **slope** of the line, and
 b is where the line crosses the y-axis
when $x=0$ (b is the **y-intercept**)





Linear Regression

- In a linear regression, where we try to formulate the relationship between variables, $y = mx + b$ becomes

$$\hat{y} = b_0 + b_1x$$

- Our goal is to predict the value of a **dependent variable** (y) based on that of an **independent variable** (x).



Linear Regression

$$\hat{y} = b_0 + b_1x$$

- How to derive b_1 and b_0 :

$$b_1 = \rho_{x,y} \frac{\sigma_y}{\sigma_x}$$

$\rho_{x,y}$ = Pearson Correlation Coefficient
 σ_x, σ_y = Standard Deviations

$$= \frac{\sum(x - \bar{x})(y - \bar{y})}{\sqrt{\sum(x - \bar{x})^2} \sqrt{\sum(y - \bar{y})^2}} \cdot \frac{\sqrt{\frac{\sum(y - \bar{y})^2}{n}}}{\sqrt{\frac{\sum(x - \bar{x})^2}{n}}} = \frac{\sum(x_i - \bar{x})(y_i - \bar{y})}{\sum(x_i - \bar{x})^2}$$



Linear Regression

$$\hat{y} = b_0 + b_1x$$

- How to derive b_1 and b_0 :

$$b_1 = \frac{\sum(x_i - \bar{x})(y_i - \bar{y})}{\sum(x_i - \bar{x})^2}$$





$$b_0 = \bar{y} - b_1\bar{x}$$



Limitations of Linear Regression

Anscombe's Quartet illustrates the pitfalls of relying on pure calculation.

Each graph results in the same calculated regression line.





Regression Exercise #1

- A manager wants to find the relationship between the number of hours that a plant is operational in a week and weekly production.





Regression Exercise #1

- Here the independent variable x is hours of operation, and the dependent variable y is production volume.





Regression Exercise #1

- The manager develops the following table:


Production Hours (x)	Production Volume (y)
34	102
35	109
39	137
42	148
43	150
47	158



Regression Exercise #1

- First, plot the data Is there a linear pattern?

Production Hours (x)	Production Volume (y)
34	102
35	109
39	137
42	148
43	150
47	158





Regression Exercise #1

- Run calculations:

$$\hat{y} = b_0 + b_1x$$
$$b_1 = \frac{\sum(x_i - \bar{x})(y_i - \bar{y})}{\sum(x_i - \bar{x})^2}$$
$$b_0 = \bar{y} - b_1\bar{x}$$

Production Hours (x)	Production Volume (y)	$(x - \bar{x})$	$(y - \bar{y})$	$(x - \bar{x})(y - \bar{y})$	$(x - \bar{x})^2$
34	102	-6	-32	192	36
35	109	-5	-25	125	25
39	137	-1	3	-3	1
42	148	2	14	28	4
43	150	3	16	48	9
47	158	7	24	168	49
\bar{x}, \bar{y}	40		Sum:	558	124
				$\sum(x - \bar{x})(y - \bar{y})$	$\sum(x - \bar{x})^2$



Regression Exercise #1

- Run calculations:

$$\begin{aligned}\hat{y} &= b_0 + b_1x \\ b_1 &= \frac{\sum(x_i - \bar{x})(y_i - \bar{y})}{\sum(x_i - \bar{x})^2} \\ b_0 &= \bar{y} - b_1\bar{x}\end{aligned}$$

Production Hours (x)	Production Volume (y)
34	102
35	109
39	137
42	148
43	150
47	158
\bar{x}, \bar{y}	40 134

$$b_1 = \frac{\sum(x_i - \bar{x})(y_i - \bar{y})}{\sum(x_i - \bar{x})^2} = \frac{558}{124} = 4.5$$

$$b_0 = \bar{y} - b_1\bar{x} = 134 - (4.5 \times 40) = -46$$

$$\hat{y} = -46 + 4.5x$$

Sum:	558	124
	$\sum(x - \bar{x})(y - \bar{y})$	$\sum(x - \bar{x})^2$



Regression Exercise #1

Based on the formula, if the manager wants to produce 125 units per week, the plant should run for:

Production Hours (x)	Production Volume (y)
34	102
35	109
39	137
42	148
43	150
47	158

$$\hat{y} = b_0 + b_1x$$

$$125 = -46 + 4.5x$$

$$x = \frac{171}{4.5} = 38 \text{ hours per week}$$



REGRESSION with Excel Data Analysis

	A	B	C						I
1	SUMMARY OUTPUT								
2									
3	Regression Statistics								
4	Multiple R	0.966875047							
5	R Square	0.934847357							
6	Adjusted R Square	0.918559196							
7	Standard Error	6.614378278							
8	Observations	6							
9									
10	ANOVA								
11		df	SS	MS	F	Significance F			
12	Regression	1	2511	2511	57.39428571	0.00162772			
13	Residual	4	175	43.75					
14	Total	5	2686						
15									
16		Coefficients	Standard Error	t Stat	P-value	Lower 95%	Upper 95%	Lower 95.0%	Upper 95.0%
17	Intercept	-46	23.91250292	-1.923679849	0.126733563	-112.3917517	20.39175167	-112.3917517	20.39175167
18	X Variable 1	4.5	0.593988704	7.575901644	0.00162772	2.85082297	6.14917703	2.85082297	6.14917703
19									

Data Analysis

Analysis Tools

Covariance
Descriptive Statistics
Exponential Smoothing
F-Test Two-Sample for Variances
Fourier Analysis
Histogram
Moving Average
Random Number Generation
Rank and Percentile
Regression

OK

Cancel

Help



LINEAR REGRESSION with Python

```
>>> from scipy.stats import linregress
>>> x = [34, 35, 39, 42, 43, 47]
>>> y = [102, 109, 137, 148, 150, 158]
>>> slope = round(linregress(x,y).slope,1)
>>> intercept = round(linregress(x,y).intercept,1)
>>> print(f'y = {intercept} + {slope}x')
y = -46.0 + 4.5x
```



Multiple Regression



Linear vs Multiple Regression

- In linear regression we have one independent variable that may relate to a dependent variable with the formula

$$\hat{y} = b_0 + b_1x$$



Linear vs Multiple Regression

- Multiple regression lets us compare several independent variables to one dependent variable at the same time.
- Each independent variable is assigned a subscript: x_1 , x_2 , x_3 etc.



Linear vs Multiple Regression

- The general formula is expanded:

linear regression

$$\hat{y} = b_0 + b_1x$$

multiple regression

$$\hat{y} = b_0 + b_1x_1 + b_2x_2 + \dots$$

- b_1 is the coefficient on x_1
- b_1 reflects the change in \hat{y} for a given change in x_1 , all else remaining constant



Linear vs Multiple Regression

- The formulas for coefficients also expand:

multiple regression

$$b_1 = \frac{\sum(x_2 - \bar{x}_2)^2 \sum(x_1 - \bar{x}_1)(y - \bar{y}) - \sum(x_1 - \bar{x}_1)(x_2 - \bar{x}_2) \sum(x_2 - \bar{x}_2)(y - \bar{y})}{\sum(x_1 - \bar{x}_1)^2 \sum(x_2 - \bar{x}_2)^2 - (\sum(x_1 - \bar{x}_1)(x_2 - \bar{x}_2))^2}$$

$$b_2 = \frac{\sum(x_1 - \bar{x}_1)^2 \sum(x_2 - \bar{x}_2)(y - \bar{y}) - \sum(x_1 - \bar{x}_1)(x_2 - \bar{x}_2) \sum(x_1 - \bar{x}_1)(y - \bar{y})}{\sum(x_1 - \bar{x}_1)^2 \sum(x_2 - \bar{x}_2)^2 - (\sum(x_1 - \bar{x}_1)(x_2 - \bar{x}_2))^2}$$

$$b_0 = \bar{y} - b_1\bar{x}_1 - b_2\bar{x}_2$$



Multiple Regression

- For example, a used car lot may want to know what variables affect net profits
- They would create a list of predictors that might correlate with profit:

price age brand
 color style






Multiple Regression

- They would want to measure the correlation of each variable to net profit
- However, some predictors might correlate with each other:

price age brand
color style





Multiple Regression

- The age of a car would have a direct impact on its sales price
- You can't adjust one without affecting the other
- This is called **multicollinearity**





Regression Exercise #2



- A pharmacy delivers medications to the surrounding community.
- Drivers can make several stops per delivery.
- The owner would like to predict the **length of time** a delivery will take based on one or two related variables.



Regression Exercise #2




- First, consider what variables may have an effect on delivery time:
 - number of stops
 - driving distance
 - outside temperature
 - gasoline prices




Regression Exercise #2




- Next, plot each variable against delivery time to see if there may be a relationship




Time vs Distance



Time vs Stops



Time vs Temperature




Time vs Gas Price



Regression Exercise #2



- Once we've chosen our variables x_1 and x_2 , we'll usually test for **multicollinearity**
- We want to know if our two independent variables are closely related to each other
- If they are, it makes sense to discard one!



Stops vs Distance

A delivery might go to one customer that lives far away, or to a group of stops close by



Regression Exercise #2



y = Delivery Time (minutes)

x_1 = Number of Stops

x_2 = Distance (miles)

y	x_1	x_2	$(y - \bar{y})$	$(x_1 - \bar{x}_1)$	$(x_1 - \bar{x}_1)^2$	$(x_2 - \bar{x}_2)$	$(x_2 - \bar{x}_2)^2$
29	1	8	-1	-1	1	2	4
31	3	4	1	1	1	-2	4
36	2	9	6	0	0	3	9
35	3	6	5	1	1	0	0
19	1	3	-11	-1	1	-3	9
\bar{y}	\bar{x}_1	\bar{x}_2			$\Sigma(x_1 - \bar{x}_1)^2$		
30	2	6			4		

$(x_1 - \bar{x}_1)(y - \bar{y})$	$(x_2 - \bar{x}_2)(y - \bar{y})$	$(x_1 - \bar{x}_1)(x_2 - \bar{x}_2)$
1	-2	-2
1	-2	-2
0	18	0
5	0	0
11	33	3
$\Sigma(x_1 - \bar{x}_1)(y - \bar{y})$	$\Sigma(x_2 - \bar{x}_2)(y - \bar{y})$	$\Sigma(x_1 - \bar{x}_1)(x_2 - \bar{x}_2)$
18	47	-1



Regression Exercise #2



y = Delivery Time (minutes)

x_1 = Number of Stops

x_2 = Distance (miles)

$$b_1 = \frac{\sum(x_2 - \bar{x}_2)^2 \sum(x_1 - \bar{x}_1)(y - \bar{y}) - \sum(x_1 - \bar{x}_1)(x_2 - \bar{x}_2) \sum(x_2 - \bar{x}_2)(y - \bar{y})}{\sum(x_1 - \bar{x}_1)^2 \sum(x_2 - \bar{x}_2)^2 - (\sum(x_1 - \bar{x}_1)(x_2 - \bar{x}_2))^2}$$

$$b_2 = \frac{\sum(x_1 - \bar{x}_1)^2 \sum(x_2 - \bar{x}_2)(y - \bar{y}) - \sum(x_1 - \bar{x}_1)(x_2 - \bar{x}_2) \sum(x_1 - \bar{x}_1)(y - \bar{y})}{\sum(x_1 - \bar{x}_1)^2 \sum(x_2 - \bar{x}_2)^2 - (\sum(x_1 - \bar{x}_1)(x_2 - \bar{x}_2))^2}$$

\bar{y}	\bar{x}_1	\bar{x}_2
30	2	6

$\sum(x_1 - \bar{x}_1)^2$
4

$\sum(x_2 - \bar{x}_2)^2$
26

$\sum(x_1 - \bar{x}_1)(y - \bar{y})$
18

$\sum(x_2 - \bar{x}_2)(y - \bar{y})$
47

$\sum(x_1 - \bar{x}_1)(x_2 - \bar{x}_2)$
-1



Regression Exercise #2



y = Delivery Time (minutes)

x_1 = Number of Stops

x_2 = Distance (miles)

$$b_1 = \frac{(26)(18) - (-1)(47)}{(4)(26) - ((-1))^2} = \frac{515}{103} = 5$$

$$b_2 = \frac{(4)(47) - (-1)(18)}{(4)(26) - ((-1))^2} = \frac{206}{103} = 2$$

\bar{y}	\bar{x}_1	\bar{x}_2
30	2	6

$\Sigma(x_1 - \bar{x}_1)^2$
4

$\Sigma(x_2 - \bar{x}_2)^2$
26

$\Sigma(x_1 - \bar{x}_1)(y - \bar{y})$
18

$\Sigma(x_2 - \bar{x}_2)(y - \bar{y})$
47

$\Sigma(x_1 - \bar{x}_1)(x_2 - \bar{x}_2)$
-1



Regression Exercise #2



y = Delivery Time (minutes)

x_1 = Number of Stops

x_2 = Distance (miles)

$$\hat{y} = 8 + 5x_1 + 2x_2$$

$$b_1 = \frac{(26)(18) - (-1)(47)}{(4)(26) - ((-1))^2} = \frac{515}{103} = 5$$

$$b_0 = \bar{y} - b_1\bar{x}_1 - b_2\bar{x}_2$$

$$= 30 - (5)(2) - (2)(6)$$

$$= 30 - 10 - 12 = 8$$

$$b_2 = \frac{(4)(47) - (-1)(18)}{(4)(26) - ((-1))^2} = \frac{206}{103} = 2$$

\bar{y}	\bar{x}_1	\bar{x}_2
30	2	6

$\Sigma(x_1 - \bar{x}_1)^2$
4

$\Sigma(x_2 - \bar{x}_2)^2$
26

$\Sigma(x_1 - \bar{x}_1)(y - \bar{y})$
18

$\Sigma(x_2 - \bar{x}_2)(y - \bar{y})$
47

$\Sigma(x_1 - \bar{x}_1)(x_2 - \bar{x}_2)$
-1



Regression Exercise #2



y = Delivery Time (minutes)

x_1 = Number of Stops

x_2 = Distance (miles)

$$\hat{y} = 8 + 5x_1 + 2x_2$$

y	x_1	x_2
29	1	8
31	3	4
36	2	9
35	3	6
19	1	3

Based on our analysis, pharmacy deliveries have a fixed time of 8 minutes, plus 5 minutes for each stop, and 2 minutes for each mile traveled



Multiple Regression in Excel

Steps are the same as linear regression, except you select a wider x-axis range

The screenshot shows the 'Regression' dialog box in Excel. The 'Input' section has 'Input Y Range' set to '\$A\$2:\$A\$6' and 'Input X Range' set to '\$B\$2:\$C\$6'. The 'Labels' checkbox is unchecked, and 'Constant is Zero' is also unchecked. The 'Confidence Level' is set to '95 %'. The 'Output options' section has 'New Worksheet Ply.' selected. The 'Residuals' section has 'Residuals', 'Standardized Residuals', 'Residual Plots', and 'Line Fit Plots' all unchecked. The 'Normal Probability' section has 'Normal Probability Plots' unchecked. The 'OK', 'Cancel', and 'Help' buttons are on the right.

Regression

Input

Input Y Range: \$A\$2:\$A\$6

Input X Range: \$B\$2:\$C\$6

☐ Labels ☐ Constant is Zero

☐ Confidence Level: 95 %

Output options

☐ Output Range:

☒ New Worksheet Ply.

☐ New Workbook

Residuals

☐ Residuals ☐ Residual Plots

☐ Standardized Residuals ☐ Line Fit Plots

Normal Probability

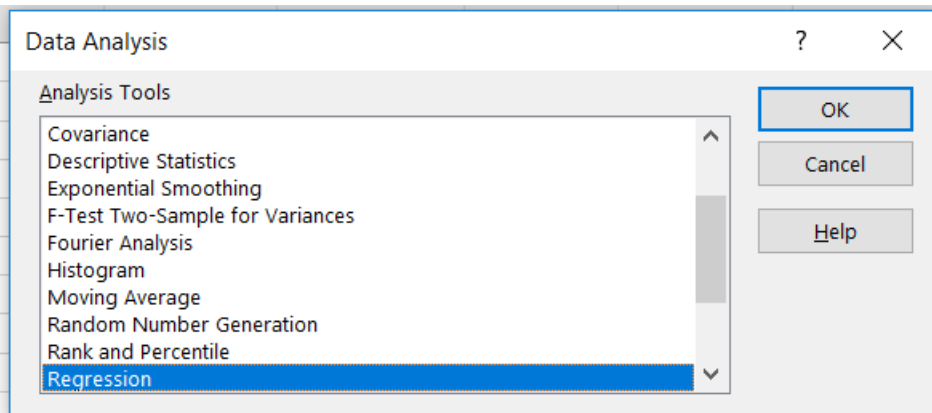
☐ Normal Probability Plots

OK Cancel Help



Multiple Regression in Excel

	A	B	C						
1	SUMMARY OUTPUT								
2									
3	Regression Statistics								
4	Multiple R	1							
5	R Square	1							
6	Adjusted R Square	1							
7	Standard Error	1.25607E-15							
8	Observations	5							
9									
10	ANOVA								
11		df	SS	MS	F	Significance F			
12	Regression	2	184	92	5.83119E+31	1.71492E-32			
13	Residual	2	3.15544E-30	1.57772E-30					
14	Total	4	184						
15									
16		Coefficients	Standard Error	t Stat	P-value	Lower 95%	Upper 95%	Lower 95.0%	Upper 95.0%
17	Intercept	8	2.11706E-15	3.77882E+15	7.00306E-32	8	8	8	8
18	X Variable 1	5	6.31078E-16	7.92295E+15	1.59304E-32	5	5	5	5
19	X Variable 2	2	2.47529E-16	8.07985E+15	1.53177E-32	2	2	2	2





MULTIPLE REGRESSION with Python

```
>>> from sklearn.linear_model import LinearRegression
>>> x1,x2 = [1,3,2,3,1], [8,4,9,6,3]
>>> y = [29,31,36,35,19]
>>> reg = LinearRegression()
>>> reg.fit(list(zip(x1,x2)), y)
>>> b1,b2 = reg.coef_[0], reg.coef_[1]
>>> b0 = reg.intercept_
>>> print(f'y = {b0:.{3}} + {b1:.{3}}x1 + {b2:.{3}}x2')
y = 8.0 + 5.0x1 + 2.0x2
```



Next Up: CHI-SQUARE ANALYSIS